

Kishor Jothimurugan

Levine 513, 3330 Walnut St, Philadelphia, PA – 19104 – USA

✉ kishor@seas.upenn.edu • 🌐 keyshor.github.io

Research Interests

My areas of interest lie at the intersection of **Formal Methods** and **Machine Learning**. In particular, I have worked on applying formal methods to improve reinforcement learning (RL), verification of neural networks (NN), and machine learning (ML) for program synthesis and analysis.

Education

University of Pennsylvania

Philadelphia, USA

Ph.D. candidate in Computer and Information Science, Current GPA 4.0/4.0

2017–2023

Thesis Topic: Specification-Guided Reinforcement Learning

Committee: Osbert Bastani, Thomas A. Henzinger, Sampath Kannan, and George J. Pappas

Advised by Prof. Rajeev Alur

Chennai Mathematical Institute

Chennai, India

B.Sc. (Honors) Mathematics and Computer Science, CGPA 9.77/10

2014–2017

Ranked among top 3 students

Publications

First (co)author in all publications; names are listed alphabetically in some papers

* equal contribution, † authors in alphabetical order

Refereed Conference Publications.....

[Festschrift 22] A Framework for Transforming Specifications in Reinforcement Learning

R. Alur, S. Bansal, O. Bastani and K. Jothimurugan[†]

Springer Festschrift in honor of Prof. Tom Henzinger (Invited), 2022

[CAV 22] Specification-Guided Learning of Nash Equilibria with High Social Welfare

K. Jothimurugan, S. Bansal, O. Bastani and R. Alur

International Conference on Computer Aided Verification, 2022

Acceptance rate: 24.4%

[NeurIPS 21] Compositional Reinforcement Learning from Logical Specifications

K. Jothimurugan, S. Bansal, O. Bastani and R. Alur

Advances in Neural Information Processing Systems, 2021

Acceptance rate: 25.7%

[EMSOFT 21] Compositional Learning and Verification of Neural Network Controllers

R. Ivanov*, K. Jothimurugan*, S. Hsu, S. Vaidya, R. Alur and O. Bastani

International Conference on Embedded Software, 2021

Acceptance rate: 25.2%

[AISTATS 21] [Abstract Value Iteration for Hierarchical Reinforcement Learning](#)

K. Jothimurugan, O. Bastani and R. Alur

International Conference on Artificial Intelligence and Statistics, 2021

Acceptance rate: 29.8%

[LICS 20] [Space-efficient Query Evaluation over Probabilistic Event Streams](#)

R. Alur, Y. Chen, K. Jothimurugan[†] and S. Khanna

Symposium on Logic in Computer Science, 2020

Acceptance rate: 39.6%

[NeurIPS 19] [A Composable Specification Language for Reinforcement Learning Tasks](#)

K. Jothimurugan, R. Alur and O. Bastani

Advances in Neural Information Processing Systems, 2019

Acceptance rate: 21.1%

[Selected \(Refereed\) Workshop Papers.....](#)

[DeepRL 22] [Robust Subtask Learning for Compositional Generalization](#)

K. Jothimurugan, S. Hsu, O. Bastani and R. Alur

Deep RL Workshop, NeurIPS, 2022

Under review at ICML 2023

[Preprints and Technical Reports.....](#)

[Policy Synthesis and Reinforcement Learning for Discounted LTL](#)

R. Alur, O. Bastani, K. Jothimurugan[†], M. Perez, F. Somenzi and A. Trivedi

Under review at CAV 2023

[Learning Algorithms for Regenerative Stopping Problems with Logistics Applications](#)

K. Jothimurugan, M. Andrews, J. Lee and L. Maggi

Nokia Bell Labs Intern Report

[Tutorials](#)

[AAAI 23] [Specification-Guided Reinforcement Learning](#)

Presenters: K. Jothimurugan, S. Bansal, R. Alur and O. Bastani

AAAI Conference on Artificial Intelligence, 2023

A comprehensive tutorial on RL algorithms for learning control policies from logical specifications

[Open Source Tools](#)

DIRL — [Github link](#)

Compositional RL algorithm for learning from temporal specifications

HIGH-NASH — [Github link](#) (CAV Artifact Evaluation Badge: Functional)

Multi-agent RL algorithm for learning Nash equilibria with high social welfare

SARL — [Github link](#)

Hierarchical RL algorithm leveraging user provided state abstractions

SPECTRL — [Github link](#)

Generating shaped rewards from temporal specifications

Teaching Experience

Guest Lecturer.....

Computer-Aided Verification, CIS 673 (Class of ~ 20) *Fall 2021*
Graduate Course, University of Pennsylvania
Topic: Techniques for Verifying Robustness of Neural Networks

Teaching Assistant.....

Principles of Embedded Systems, CIS 540 (Class of ~ 20) *Spring 2019*
Graduate Course, University of Pennsylvania

Automata, Computability and Complexity, CIS 262 (Class of ~ 150) *Fall 2018*
Undergraduate Course, University of Pennsylvania

Discrete Mathematics (Class of ~ 30) *Spring 2017*
Undergraduate Course, Chennai Mathematical Institute

Design and Analysis of Algorithms *Fall 2016*
Online Course offered by Indian Institute of Technology, Madras

Mentoring and Outreach

Graduate Student Mentoring 2020–2022
Steve Hsu, Masters Student, University of Pennsylvania
Topic: Compositional reinforcement learning for multi-task generalization
Co-authored two papers

Penn CIS Mentorship Program 2022
Participated as a mentor
Department-wide effort to connect early career PhD students with senior PhD students
Attracted participation from 34 first-year students

Awards

CTL Teaching Certificate *Spring 2022*
Awarded by Center for Teaching and Learning, University of Pennsylvania

CMI Undergraduate Scholarship *2014–2017*
Awarded by CMI to undergraduate students for excellence in academics

Research Presentations

Invited Talks.....

University of Illinois Urbana-Champaign *Spring 2023*
Title: Specification-Guided Reinforcement Learning

University of Virginia *Spring 2023*
Title: Specification-Guided Reinforcement Learning

NYU Formal Methods Group Fall 2022
Title: Specification-Guided Reinforcement Learning

Microsoft PROSE Group Fall 2022
Title: Specification-Guided Reinforcement Learning

Thomas Henzinger's Group, IST Austria Fall 2021
Title: Reinforcement Learning from Logical Specifications

Simons Institute, UC Berkeley Spring 2021
Workshop on Games and Equilibria in System Design and Analysis
Title: Abstract Value Iteration for Hierarchical Reinforcement Learning

Presentations and Posters (from publications).....

Conferences: NeurIPS 2019 and 2021, LICS 2020, AISTATS 2021, EMSOFT 2021, CAV 2022.

Workshops: DeepRL @ NeurIPS 2020 and 2022, SafeRL @ NeurIPS 2021, SYNT @ CAV 2022.

Internships

Amazon Web Services Summer 2022
Applied Science Intern, AI Labs
Incorporating program semantics in transformer-based code generation models
Mentors: Nathan Fulton, Siddhartha Jain and Baishakhi Ray (Columbia University)

Nokia Bell Labs Summer 2020
Research Intern, Data and AI Lab
An application of deep RL to regenerative stopping problems
Mentors: Matthew Andrews, Jeongran Lee and Lorenzo Maggi

Amazon Web Services Summer 2019
Software Development Engineering Intern, Automated Reasoning Group
Using machine learning to improve usability of taint analysis
Mentors: Andrew Gacek and Lee Pike

ENS Paris-Saclay Summer 2017
Research Intern, Formal Methods Group
Models for distributed reactive synthesis
Mentor: Dietmar Berwanger

Review Service

Conferences: NeurIPS 2022, AAAI 2023, AISTATS 2023, ESOP 2023, L4DC 2023, ICML 2023, CAV 2023, NeurIPS 2023.

Journals: IEEE TCAD.